

Learnable Drift

Temporal ACS and covariate/feature shifts contain stable signal; anchored policy improves mean F1.

Label-Flip Stress

Subgroup label flips inject contradictions; static deployment outperforms retraining.

Ablation separates useful adaptation from noise chasing

Policy Gate

Persistence and cooldown reduce unnecessary adaptation compared with immediate retraining.

Core Finding

Drift detection is not enough: retraining helps only when new labels describe a learnable regime.